

cluster-over-sampling: A package for clustering-based oversampling

Georgios Douzas*, Fernando Bacao

NOVA Information Management School, Universidade Nova de Lisboa

Abstract

Learning from imbalanced data is a common and challenging problem in supervised learning. Standard classifiers are designed to handle balanced class distributions. While different strategies exist to tackle this problem, methods that generate artificial data to achieve a balanced class distribution, called oversampling algorithms, are more versatile than modifications to the classification algorithms. SMOTE algorithm, the most popular oversampler, as well as any other oversampling method based on it, generates synthetic samples along line segments that join minority class instances. SMOTE addresses only the issue of between-classes imbalance. On the other hand, by clustering the input space and applying any oversampling algorithm for each resulting cluster with appropriate resampling ratio, the within-classes imbalanced issue can be addressed. This approach, implemented in the **cluster-over-sampling** Python open source project, has been shown in multiple publications, using a variety of datasets, to outperform other standard oversamplers. In this paper we describe **cluster-over-sampling** in detail and make it available to the machine learning community. An important point is that the implementation integrates effortlessly with the Scikit-Learn ecosystem. Therefore, machine learning researchers and practitioners can integrate it directly to any pre-existing work.

Keywords: Machine learning, Classification, Imbalanced learning, Oversampling, Clustering

*Postal Address: NOVA Information Management School, Campus de Campolide, 1070-312 Lisboa, Portugal, Telephone: +351 21 382 8610

Email addresses: gdouzas@novaims.unl.pt (Georgios Douzas),
bacao@novaims.unl.pt (Fernando Bacao)

Code metadata	
Current code version	v0.1.1
Permanent link to code/repository used for this code version	https://github.com/AlgoWit/cluster-over-sampling
Legal Code License	MIT
Code versioning system used	git
Software code languages, tools, and services used	Python, Travis CI, AppVeyor, Read the Docs, Codecov, CircleCI, zenodo, Anaconda Cloud
Compilation requirements, operating environments & dependencies	Linux, Mac OS, Windows
If available Link to developer documentation/manual	https://cluster-over-sampling.readthedocs.io/
Support email for questions	georgios.douzas@gmail.com

Table 1: Code metadata

1. Motivation and significance

1.1. Introduction

The imbalanced learning problem describes the case where in a machine learning classification task, using datasets with binary or multi-class targets, one of the classes, called the majority class, has a significantly higher number of samples compared to the remaining classes, called the minority class(es) [1]. Learning from imbalanced data is a non-trivial problem for both academic researchers and industry practitioners that can be frequently found in multiple domains such as chemical and biochemical engineering, financial management, information technology, security, business, agriculture or emergency management [2].

A bias towards the majority class is induced when imbalanced data are used to train standard machine learning algorithms. This results in low classification accuracy, especially for the minority class(es), when the classifier is evaluated on unseen data. An important measure for the degree of data imbalance is the Imbalance Ratio (IR), defined as the ratio between the number of samples of the majority class and each of the minority classes. Using a rare disease detection task as an example, with 1% of positive cases corresponding to an $IR = \frac{0.99}{0.01} = 99$, a trivial classifier that always labels a person as healthy will score a classification accuracy of 99%. However in this case, all positive cases remain undetected. The observed values of IR are

often between 100 and 100.000 [3], [4]. Figure 1 presents an example of imbalanced data in two dimensions as well as the decision boundary identified by a typical classifier when they are used as a training set.

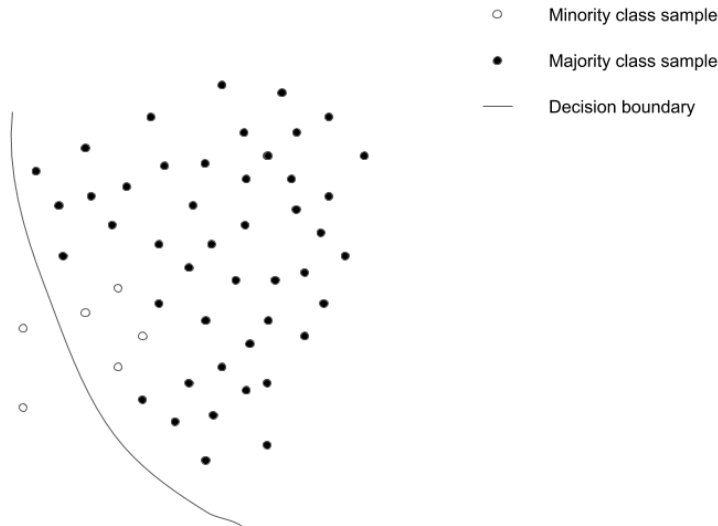


Figure 1: Imbalanced data in two dimensions. The decision boundary of a typical classifier shows a bias towards the majority class.

1.2. Oversampling algorithms

Various approaches have been proposed to improve classification results when the training data are imbalanced, a case also known as between-class imbalance. The most general approach, called oversampling, is the generation of artificial data for the minority class(es) [5]. Synthetic Minority Oversampling Technique (SMOTE) [3] was the first non-trivial oversampler proposed and remains the most popular one. Although SMOTE has been shown to be effective for generating artificial data, it also has some drawbacks [6]. In order to improve the quality of the artificial data many variants of SMOTE have been proposed. Nevertheless, they utilize the SMOTE data generation mechanism, which consists of a linear interpolation between minority class samples to generate synthetic instances as shown in figure 2.




Figure 2: Visual representation of the SMOTE data generation mechanism.

37 A Python implementation of SMOTE and several of its variants is avail-
 38 able in the Imbalanced-Learn [7] library, which is fully compatible with the
 39 popular machine learning toolbox Scikit-Learn [8].

40 1.3. Clustering-based oversampling

41 In addition to between-class imbalance, within-class imbalance refers to
 42 the case where areas of sparse and dense minority class instances exist. As
 43 the first step of generating synthetic samples, the SMOTE data generation
 44 mechanism selects randomly, with uniform probability, minority class in-
 45 stances. Consequently, dense minority class areas have a high probability
 46 of being inflated further, while the sparsely populated are likely to remain
 47 sparse. This allows to combat between-class imbalance, while the issue of
 48 within-class imbalance is ignored [9].

49 On the other hand, clustering-based oversampling, as presented in [10]
 50 and [11], aims to deal with both between-class and within-class imbalance
 51 problems. Initially a clustering algorithm is applied to the input space. The
 52 resulting clusters allow to identify sparse and dense minority class(es) areas.
 53 A small IR, relatively to a threshold, of a particular cluster is used as an
 54 indicator that it can be safely used as a data generation area, i.e. noise
 55 generation is avoided. Furthermore, sparse minority clusters are assigned
 56 more synthetic samples, which alleviates within-class imbalance.

57 Specific realizations of the above approach are SOMO [10] and KMeans-
 58 SMOTE [11] algorithms. Empirical studies have shown that both algorithms

59 outperform SMOTE and its variants across multiple imbalanced datasets,
60 classifiers and evaluation metrics. In this paper, we present a generic Python
61 implementation of clustering-based oversampling, in the sense that any com-
62 bination of a Scikit-Learn compatible clusterer and Imbalanced-Learn com-
63 batible oversampler can be selected to produce an algorithm that identifies
64 clusters on the input space and applies oversampling on each one of them.
65 In section 2, the software description is given while section 3 provides a
66 demonstrative example of its functionalities.

67 2. Software description

68 The `cluster-over-sampling` software project is written in Python 3.7.
69 It contains an object-oriented implementation of the clustering-based over-
70 sampling procedure as well as detailed online documentation. The imple-
71 mentation provides an API that is compatible with Imbalanced-Learn and
72 Scikit-Learn libraries. Therefore, standard machine learning functionalities
73 are supported while the generated clustering-based oversampling algorithm,
74 composed by a clusterer and an oversampler, contains the initial oversampler
75 functionality as a special case.

76 2.1. Software architecture

77 The `cluster-over-sampling` project contains the Python package `clover`.
78 The main modules of `clover` are called `distribution` and `over_sampling`.
79 The `distribution` module implements the functionality related to the distri-
80 bution of the generated samples to the identified clusters, while `over_sampling`
81 implements the functionality related to the generation of artificial samples.
82 Both of them are presented in detail below.

83 2.1.1. Module *distribution*

84 The module `distribution` contains the files `base.py` and `density.py`.
85 The former provides the implementation of the `BaseDistributor` class, the
86 base class for distributors, while the later includes the `DensityDistributor`
87 class, a generalization of the density based distributor presented in [10] and
88 [11], that inherits from `BaseDistributor`. Following the Scikit-Learn API,
89 `BaseDistributor` includes the public methods `fit` and `fit_distribute`.
90 The `fit_distribute` method calls the `fit` method and returns two Python
91 dictionaries that describe the distribution of generated samples inside each
92 cluster and between clusters, respectively. Specifically, the `fit` method calcu-
93 lates various statistics related to the distribution process, while it calls `_fit`
94 method to calculate the actual intra-cluster and inter-cluster distributions.
95 This is achieved by invoking the `_intra_distribute` and `_inter_distribute`

96 methods. The `BaseDistributor` class provides a trivial implementation of
 97 them, that should be overwritten when a realization of a distributor class is
 98 considered. Consequently, `DensityDistributor` overwrites both methods as
 99 well as the `_fit` method. The later calls the methods `_identify_filtered_clusters`
 100 and `_calculate_clusters_density` that identify the clusters used for data
 101 generation and calculate their density, respectively. Subsection 2.2 pro-
 102 vides a detailed description of the initialization and functionality of the
 103 `DensityDistributor` class. Figure 3 shows a visual representation of the
 104 above classes and functions hierarchy.

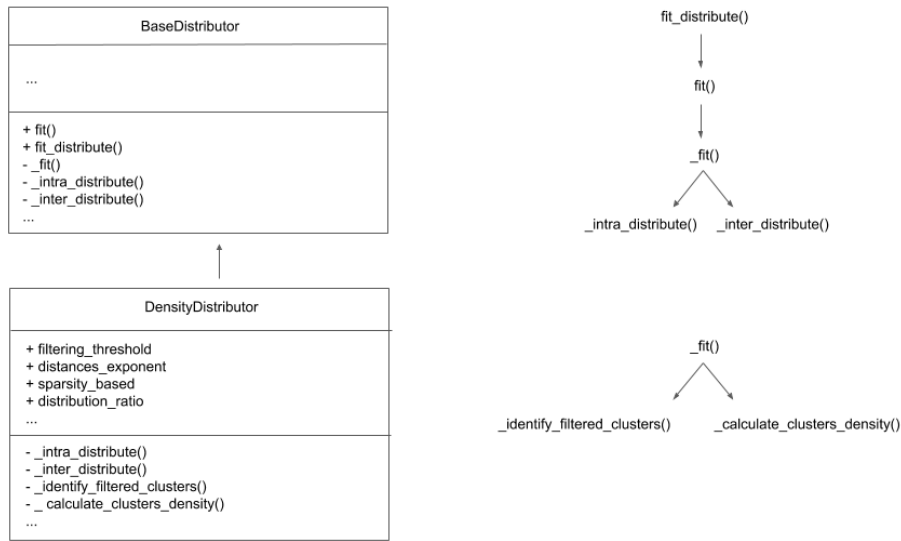


Figure 3: UML `BaseDistributor` and `DensityDistributor` class diagrams and callgraphs of main classes and methods.

105 2.1.2. Module *oversampling*

106 The module `over_sampling` contains the files `base.py` and `monkey_patching.py`.
 107 The former provides the implementation of the `BaseClusterOverSampler`
 108 class, an extension of the Imbalanced-Learn's `BaseOverSampler` class, while
 109 the later enhances the main oversamplers provided by Imbalanced-Learn
 110 with the functionality required by clustering-based oversampling. The ini-
 111 tializer of `BaseClusterOverSampler`, compared to `BaseOverSampler`, in-
 112 cludes the extra parameters `clusterer` and `distributor`. Also following
 113 the Imbalanced-Learn API, `BaseClusterOverSampler` includes the public
 114 methods `fit` and `fit_resample`. It also inherits from `BaseOverSampler`, the
 115 base class of oversamplers that is implemented in Imbalanced-Learn. The `fit`
 116 method calculates various statistics related to the resampling process, while

117 the `fit_resample` method returns an enhanced version of the input data
 118 by appending the artificially generated samples. Specifically, `fit_resample`
 119 calls the `_fit_resample` method that in turn calls the `_intra_sample` and
 120 `_inter_sample` methods to generate the intra-cluster and inter-cluster arti-
 121 ficial samples, respectively. This is achieved by invoking the `_fit_resample_cluster`
 122 method that implements the data generation mechanism. Therefore every
 123 oversampler that inherits from the `BaseClusterOverSampler` class should
 124 overwrite `_fit_resample_cluster`, providing a concrete implementation of
 125 the oversampling process. This has been done for the main oversamplers of
 126 Imbalanced-Learn through the `monkey_patching` module. Specifically, for
 127 each oversampler the `_fit_resample_cluster` method has been set equal to
 128 the `_fit_resample` method. Subsection 2.2 provides a detailed description of
 129 the initialization and functionality of the various oversamplers, enhanced by
 130 the clustering process. Figure 4 shows a visual representation of the above
 131 classes and functions hierarchy.

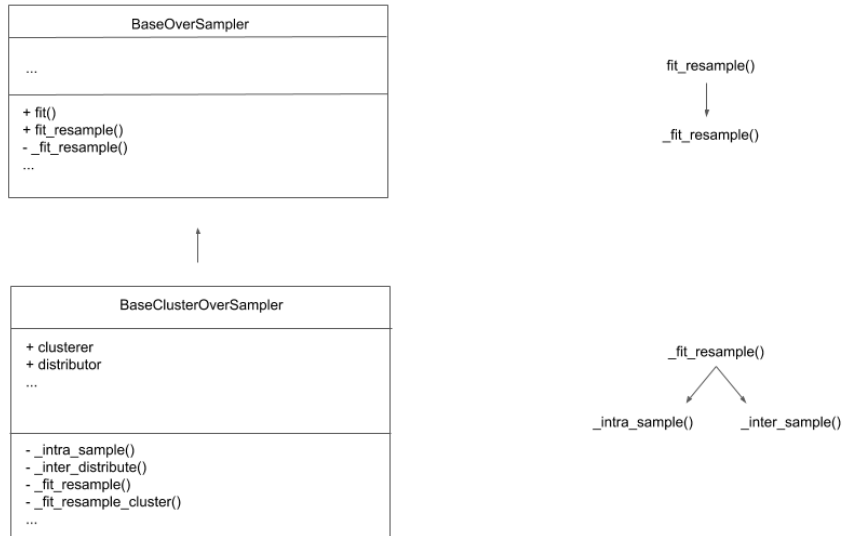


Figure 4: UML `BaseOverSampler` and `BaseClusterOverSampler` class diagrams and call-graphs of main classes and methods.

132 2.2. Software Functionalities

133 As it was mentioned in subsection 2.1.2, `cluster-over-sampling` ex-
 134 tends Imbalanced-Learn’s functionality by clustering the input space be-
 135 fore oversampling is applied. This is achieved through the implementation
 136 of the `BaseClusterOverSampler` class, an extension of Imbalanced-Learn’s
 137 `BaseOverSampler` class. Oversamplers that inherit from `BaseClusterOverSampler`,

138 compared to oversamplers inheriting from `BaseOverSampler`, require two ad-
 139 ditional initialization parameters: `clusterer` and `distributor`. Their de-
 140 fault values are for both parameters equal to `None`, a case that corresponds
 141 to the usual oversampling procedure i.e. no clustering applied to the input
 142 space. On the other hand if the parameter `clusterer` is equal to any Scikit-
 143 Learn compatible clustering algorithm then clustering of the input space is
 144 initially applied, followed by oversampling in each cluster with the distribu-
 145 tion of generated samples given by the `distributor` parameter. The default
 146 `distributor` value is an instance of `DensityDistributor` class as described
 147 in subsection 2.1.1. The initializer of `DensityDistributor` includes the fol-
 148 lowing parameters: `filtering_threshold`, `distances_exponent`, `sparsity_based`
 149 and `distribution_factor`. The first parameter is used to identify the fil-
 150 tered clusters, i.e. clusters of samples that are included in the data generation
 151 process. The second parameter modifies the density calculation of the filtered
 152 clusters by increasing the effect of euclidean distances between samples. The
 153 third parameter selects whether generated samples are assigned to filtered
 154 clusters inversly proportional to their density. Finally the last parameter ad-
 155 justs the intra-cluster to inter-cluster proportion of generated samples, while
 156 it applies only to clusterers that support a neighborhood structure. Once the
 157 `DensityDistributor` object is initialized with a specific parametrization, it
 158 can be used to distribute the generated samples to the clusters identified by
 159 any clustering algorithm.

160 Resampling is achieved by using the two main methods of `fit` and `fit_resample`
 161 of any oversampler inheriting from `BaseClusterOverSampler`. More specif-
 162 ically, both of them take as input parameters the input matrix `X` and target
 163 labels `y`. Following the Scikit-Learn API, both `X`, `y` are array-like objects of
 164 appropriate shape. The first method computes various statistics which are
 165 used to resample `X`, while the second method does the same but additionally
 166 returns a resampled version of `X` and `y`.

167 The `cluster-over-sampling` project has been designed to integrate with
 168 the Imbalanced-Learn toolbox and the Scikit-Learn ecosystem. Therefore
 169 any oversampler that inherits from `BaseClusterOverSampler` can be used
 170 in a machine learning pipeline, through Imbalanced-Learn’s class `Pipeline`,
 171 that automatically combines `samplers`, `transformers` and `estimators`. The
 172 next section provides examples of the above functionalities.

173 3. Illustrative examples

174 3.1. Basic example

175 An example of resampling an imbalanced dataset using the `fit_resample`
 176 method is presented in Listing 1. Initially, a binary-class imbalanced dataset

177 is generated. Next, a **SMOTE** oversampler is initialized using **KMeans** as a
 178 clusterer. This effectively corresponds to the KMeans-SMOTE algorithm as
 179 presented in [11]. Finally, the oversampler's **fit_resample** method is used to
 180 resample the data. Printing the class distribution before and after resampling
 181 confirms that the resampled data **X_res**, **y_res** are perfectly balanced. **X_res**,
 182 **y_res** can be used as training data for any classifier in the place of **X**, **y**.

Listing 1: Resampling of imbalanced data using the **fit_resample** method of KMeans-SMOTE oversampling algorithm.

```

183 # Import classes and functions.
184 from collections import Counter
185 from clover.over_sampling import SMOTE
186 from sklearn.cluster import KMeans
187 from sklearn.datasets import make_classification
188
189 # Generate an imbalanced a binary class dataset.
190 X, y = make_classification(
191     random_state=23,
192     n_classes=2,
193     n_features=5,
194     n_samples=1000,
195     weights=[0.8, 0.2]
196 )
197
198 # Create KMeans-SMOTE object with default hyperparameters.
199 kmeans_smote = SMOTE(clusterer=KMeans(random_state=4), random_state=10
200
201 # Resample the imbalanced dataset.
202 X_res, y_res = kmeans_smote.fit_resample(X, y)
203
204 # Print number of samples per class for initial and resampled data.
205 init_count = list(Counter(y).values())
206 resampled_count = list(Counter(y_res).values())
207
208 print(f'Initial_class_distribution:_{init_count}.')
209 # Initial class distribution: [792, 208].
210
211 print(f'Resampled_class_distribution:_{resampled_count}.')
212 # Resampled class distribution: [792, 792].

```

213 3.2. Machine learning pipeline

214 As mentioned before, any clustering-based oversampler can be used as a
 215 part of a machine learning pipeline. Listing 2 presents a pipeline composed
 216 by the combination of Borderline SMOTE oversampler and hierarchical clus-
 217 tering, a PCA tranformation and a decision tree classifier. The pipeline is
 218 trained on multi-class imbalanced data and evaluated on a hold-out set. The
 219 user applies the process in a simple way while the internal details of the
 220 calculations are hidden.

Listing 2: Training and evaluation of a machine learning pipeline that contains the AgglomerativeClustering-BorderlineSMOTE algorithm.

```

221 # Import classes and functions.
222 from clover.over_sampling import BorderlineSMOTE
223 from sklearn.datasets import make_classification
224 from sklearn.decomposition import PCA
225 from sklearn.tree import DecisionTreeClassifier
226 from sklearn.model_selection import train_test_split
227 from sklearn.metrics import f1_score
228 from sklearn.cluster import AgglomerativeClustering
229 from imblearn.pipeline import make_pipeline
230
231 # Generate an imbalanced multi-class dataset.
232 X, y = make_classification(
233     random_state=23,
234     n_classes=3,
235     n_informative=10,
236     n_samples=500,
237     weights=[0.8, 0.1, 0.1]
238 )
239
240 # Split the data to training and hold-out sets.
241 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state
242
243 # Create the pipeline's objects with default hyperparameters.
244 hclusterer_bsmote = BorderlineSMOTE(clusterer=AgglomerativeClustering(
245     pca = PCA()
246     clf = DecisionTreeClassifier(random_state=3)
247
248 # Create the pipeline.
249 pip = make_pipeline(hclusterer_bsmote, pca, clf)
250

```

```

251 # Fit the pipeline to the training set.
252 pip.fit(X_train, y_train)
253
254 # Evaluate the pipeline on the hold-out set using the F-score.
255 test_score = f1_score(y_test, pip.predict(X_test), average='micro')
256
257 print(f'F-score on hold-out set: {test_score:.2f}')
258 # F-score on hold-out set: 0.78.

```

259 4. Impact and conclusions

260 Classification of imbalanced datasets is a challenging task for standard
 261 machine learning algorithms. In addition to between-class imbalance, within-
 262 class imbalance refers to the case where areas of sparse and dense minority
 263 class instances exist. Clustering-based oversampling, aims to deal with both
 264 between-class and within-class imbalance problems.

265 The **cluster-over-sampling** project provides the only Python imple-
 266 mentation, to the best of our knowledge, that provides a generic way to
 267 construct any clustering-based oversampler. A significant advantage of this
 268 implementation is that it is built on top of the Scikit-Learn’s ecosystem and
 269 therefore it can be easily used in typical machine learning workflows. Also,
 270 the public API of any clustering-based oversampler is an extension of the
 271 one provided in Imbalanced-Learn. This means that users of Imbalanced-
 272 Learn and Scikit-Learn, that apply oversampling on imbalanced data, can
 273 integrate the **cluster-based-oversampler** package in their existing work
 274 in a straightforward manner.

275 5. Conflict of interest

276 We wish to confirm that there are no known conflicts of interest associated
 277 with this publication and there has been no significant financial support for
 278 this work that could have influenced its outcome.

279 References

- 280 [1] N. V. Chawla, A. Lazarevic, L. Hall, K. Boyer, SMOTEBoost: im-
 281 proving prediction of the minority class in boosting, Principles of
 282 Knowledge Discovery in Databases, PKDD-2003 (2003) 107–119doi:
 283 10.1007/b13634.
 284 URL [http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.](http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.80.1499)
 285 [1.80.1499](http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.80.1499)

- [2] G. Haixiang, L. Yijing, J. Shang, G. Mingyun, H. Yuanyue, G. Bing, Learning from class-imbalanced data: Review of methods and applications, *Expert Systems with Applications* 73 (2017) 220–239. doi:10.1016/j.eswa.2016.12.035. URL <https://doi.org/10.1016/j.eswa.2016.12.035>
- [3] N. V. Chawla, K. W. Bowyer, L. O. Hall, W. P. Kegelmeyer, SMOTE: Synthetic minority over-sampling technique, *Journal of Artificial Intelligence Research* 16 (2002) 321–357. arXiv:1106.1813, doi:10.1613/jair.953.
- [4] S. Barua, M. M. Islam, X. Yao, K. Murase, MWMOTE - Majority weighted minority oversampling technique for imbalanced data set learning, *IEEE Transactions on Knowledge and Data Engineering* 26 (2) (2014) 405–425. doi:10.1109/TKDE.2012.232.
- [5] A. Fernández, V. López, M. Galar, M. J. del Jesus, F. Herrera, Analysing the classification of imbalanced data-sets with multiple classes: Binarization techniques and ad-hoc approaches, *Knowledge-Based Systems* 42 (2013) 97–110. doi:http://dx.doi.org/10.1016/j.knosys.2013.01.018. URL <http://www.sciencedirect.com/science/article/pii/S0950705113000300>
- [6] H. He, E. A. Garcia, Learning from Imbalanced Data, *IEEE Transactions on Knowledge and Data Engineering* 21 (9) (2009) 1263–1284. arXiv:arXiv:1011.1669v3, doi:10.1109/TKDE.2008.239.
- [7] G. Lemaitre, F. Nogueira, C. K. Aridas, Imbalanced-learn: A Python Toolbox to Tackle the Curse of Imbalanced Datasets in Machine Learning, *Journal of Machine Learning Research* 18 (2016) 1–5. arXiv:1609.06570, doi:http://www.jmlr.org/papers/volume18/16-365/16-365.pdf. URL <http://arxiv.org/abs/1609.06570>
- [8] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, E. Duchesnay, *Scikit-learn: Machine learning in Python*, Vol. 12, 2011. arXiv:arXiv:1201.0490v2, doi:10.1007/s13398-014-0173-7.2. URL <http://dl.acm.org/citation.cfm?id=2078195>

- 321 [9] R. C. Prati, G. Batista, M. C. Monard, Learning with class skews
322 and small disjuncts, in: SBIA, 2004, pp. 296–306. doi:10.1007/
323 978-3-540-28645-5_30.
- 324 [10] G. Douzas, F. Bacao, Self-organizing map oversampling (somo) for im-
325 balanced data set learning, Expert Systems with Applications 82 (2017)
326 40 – 52. doi:https://doi.org/10.1016/j.eswa.2017.03.073.
327 URL [http://www.sciencedirect.com/science/article/pii/
328 S0957417417302324](http://www.sciencedirect.com/science/article/pii/S0957417417302324)
- 329 [11] G. Douzas, F. Bacao, F. Last, Improving imbalanced learn-
330 ing through a heuristic oversampling method based on k-
331 means and smote, Information Sciences 465 (2018) 1 – 20.
332 doi:https://doi.org/10.1016/j.ins.2018.06.056.
333 URL [http://www.sciencedirect.com/science/article/pii/
334 S0020025518304997](http://www.sciencedirect.com/science/article/pii/S0020025518304997)