

cluster-over-sampling: A package for clustering based oversampling

Georgios Douzas, Fernando Bacao*

NOVA Information Management School, Universidade Nova de Lisboa

Abstract

Learning from class-imbalanced data is a common and challenging problem in supervised learning. Standard classification algorithms are designed to handle balanced class distributions. While different strategies exist to tackle this problem, methods that generate artificial data to achieve a balanced class distribution, called oversampling algorithms, are more versatile than modifications to the classification algorithms. SMOTE algorithm, the most popular oversampler, as well as any other oversampling method based on it, generates synthetic samples along line segments that join minority class instances. SMOTE addresses only the issue of between-classes imbalance. On the other hand, by clustering the input space and applying any oversampling algorithm for each resulting cluster with appropriate resampling ratio, the within-classes imbalanced issue can be addressed. This approach, implemented in the cluster-over-sampling Python open source project, has been shown in various publications to outperform other standard oversamplers in a large number of datasets. In this paper we describe cluster-over-sampling in detail and make it available to the machine learning community. An important point is that the implementation integrates effortlessly with the Scikit-Learn ecosystem. Therefore, machine learning researchers and practitioners can integrate it directly to any pre-existing work.

Keywords: Machine learning, Classification, Imbalanced learning, Oversampling

*Postal Address: NOVA Information Management School, Campus de Campolide, 1070-312 Lisboa, Portugal, Telephone: +351 21 382 8610

Email addresses: `gdouzas@novaims.unl.pt` (Georgios Douzas),
`bacao@novaims.unl.pt` (Fernando Bacao)

Code metadata	
Current code version	v0.1.1
Permanent link to code/repository used for this code version	https://github.com/AlgoWit/cluster-over-sampling
Legal Code License	MIT
Code versioning system used	git
Software code languages, tools, and services used	Python, Travis CI, AppVeyor, Read the Docs, Codecov, CircleCI, zenodo, Anaconda Cloud
Compilation requirements, operating environments & dependencies	Linux, Mac OS, Windows
If available Link to developer documentation/manual	https://cluster-over-sampling.readthedocs.io/
Support email for questions	georgios.douzas@gmail.com

Table 1: Code metadata

1. Motivation and significance

1.1. Introduction

The imbalanced learning problem is defined as a machine learning classification task using datasets with binary or multi-class targets where one of the classes, called the majority class, outnumbers significantly the remaining classes, called the minority class(es) [1]. Learning from imbalanced data is a frequent and non-trivial problem for academic researchers and industry practitioners alike. The imbalance learning problem can be found in multiple domains such as chemical and biochemical engineering, financial management, information technology, security, business, agriculture or emergency management [2].

The imbalanced learning problem describes the case where in a machine learning classification task using datasets with binary or multi-class targets, one of the classes, called the majority class, has a significantly higher number of samples compared to the remaining classes, called the minority class(es) [1]. Learning from imbalanced data is a non-trivial problem for both academic researchers and industry practitioners. Additionally, imbalanced data can be frequently found in multiple domains such as chemical and biochemical engineering, financial management, information technology, security, business, agriculture or emergency management [2].

A bias towards the majority class is induced when imbalanced data are used to train standard machine learning algorithms. This results in low clas-

23 sification accuracy, especially for the minority class(es), when the classifier
 24 is evaluated on unseen data. An important measure for the degree of data
 25 imbalance is the Imbalance Ratio (IR), defined as the ratio between the
 26 number of samples of the majority class and each of the minority classes.
 27 Using a rare disease detection task as an example, with 1% of positive cases
 28 corresponding to an $IR = \frac{0.99}{0.01} = 99$, a trivial classifier that always labels a
 29 person as healthy will score a classification accuracy of 99%. However in this
 30 case, all positive cases remain undetected. The observed values of IR are
 31 often between 100 and 100.000 [3], [4]. Figure 1 presents an example of im-
 32 balanced data in two dimensions as well as the decision boundary identified
 33 by a typical classifier when they are used as a training set.




Figure 1: Imbalanced data in two dimensions. The decision boundary of a typical classifier shows a bias towards the majority class.

34 1.2. *Oversampling algorithms*

35 Various approaches have been proposed to improve classification results
 36 when the training data are imbalanced, a case also known as between-class
 37 imbalance. The most general approach, called oversampling, is the generation
 38 of artificial data for the minority class(es) [5]. Synthetic Minority Oversam-
 39 pling Technique (SMOTE) [3] was the first non-trivial oversampler proposed
 40 and remains the most popular one. Although SMOTE has been shown to
 41 be effective for generating artificial data, it also has some drawbacks [6]. In
 42 order to improve the quality of the artificial data many variants of SMOTE
 43 have been proposed. Nevertheless, they utilize the SMOTE data generation

44 mechanism, which consists of a linear interpolation between minority class
 45 samples to generate synthetic instances as shown in figure 2.




Figure 2: Visual representation of the SMOTE data generation mechanism.

46 A Python implementation of SMOTE and several of its variants is avail-
 47 able in the Imbalanced-Learn [7] library, which is fully compatible with the
 48 popular machine learning toolbox Scikit-Learn [8].

49 1.3. Clustering based oversampling

50 In addition to between-class imbalance, within-class imbalance refers to
 51 the case where areas of sparse and dense minority class instances exist. As a
 52 first step of generating synthetic samples, the SMOTE data generation mech-
 53 anism selects randomly, with uniform probability, minority class instances.
 54 Consequently, dense minority class areas have a high probability of being in-
 55 flated further, while the sparsely populated are likely to remain sparse. This
 56 allows to combat between-class imbalance, while the issue of within-class
 57 imbalance is ignored [9].

58 On the other hand, clustering based oversampling, as presented in [10]
 59 and [11], aims to deal with both between-class and within-class imbalance
 60 problems. Initially a clustering algorithm is applied to the input space. The
 61 resulting clusters allow to identify sparse and dense minority class(es) areas.
 62 A small IR, compared to a threshold, of a particular cluster is used as an
 63 indicator that it can be safely used as a safe data generation area, i.e. noise

64 generation is avoided. Furthermore, sparse minority clusters are assigned
65 more synthetic samples, which alleviates within-class imbalance.

66 Specific realizations of the above approach are SOMO [10] and KMeans-
67 SMOTE [11] algorithms. Empirical studies have shown that both algorithms
68 outperform SMOTE and its variants across multiple imbalanced datasets,
69 classifiers and evaluation metrics. In this paper, we present a generic Python
70 implementation of clustering based oversampling, in the sense that any com-
71 bination of a Scikit-Learn compatible clusterer and Imbalanced-Learn com-
72 batible oversampler can be selected to produce an algorithm that identifies
73 clusters on the input space and apply oversampling on each one of them. In
74 section 2, the software description is given while section 3 provides a demon-
75 strative example of its functionalities.

76 2. Software description

77 The `cluster-over-sampling` software project is written in Python 3.7.
78 It contains an object-oriented implementation of the cluster based oversam-
79 pling procedure as well as detailed online documentation. The implementa-
80 tion provides an API that is compatible with Imbalanced-Learn and Scikit-
81 Learn libraries. Therefore, standard machine learning functionalities are sup-
82 ported while the generated clustering based oversampling algorithm includes
83 any selected oversampler as a special case.

84 2.1. Software Architecture

85 The `cluster-over-sampling` project contains the Python package `clover`.
86 The main modules of `clover` are called `distribution` and `over_sampling`.

87 The module `distribution` implements the functionality related to the
88 distribution of the generated samples to the identified clusters. It con-
89 tains the files `base.py` and `density.py`. The former provides the imple-
90 mentation of the `BaseDistributor` class, the base class for distributors,
91 while the later includes the `DensityDistributor` class, a generalization of
92 the density based distributor presented in [10] and [11], that inherits from
93 `BaseDistributor`. Following the Scikit-Learn API, `BaseDistributor` in-
94 cludes the public methods `fit` and `fit_distribute`. The `fit_distribute`
95 method calls the `fit` method and returns two Python dictionaries that
96 describe the distribution of generated samples inside each cluster and be-
97 tween clusters, respectively. Particularly, the `fit` method calculates vari-
98 ous statistics related to the distribution process, while it calls `_fit` method
99 to calculate the actual intra-cluster and inter-cluster distributions. This is
100 achieved by invoking the `_intra_distribute` and `_inter_distribute` meth-
101 ods. The `BaseDistributor` class provides a trivial implementation of them,

that should be overwritten when a realization of a distributor class is considered. Therefore, `DensityDistributor` overwrites both methods as well as the `_fit` method. The later calls the methods `_identify_filtered_clusters` and `_calculate_clusters_density` that identify the clusters used for data generation and calculate their density, respectively. Subsection 2.2 provides a detailed description of the initialization and functionality of the `DensityDistributor` class. Figure ?? shows a visual representation of the above classes and functions hierarchy.

The initialization of a `GeometricSMOTE` instance includes G-SMOTE's hyperparameters that control the generation of synthetic data. Additionally, `GeometricSMOTE` inherits from the `BaseOverSampler` class of Imbalanced-Learn library. Therefore, an instance of `GeometricSMOTE` class provides the `fit` and `fit_resample` methods, the two main methods for resampling as explained in subsection 2.2. This is achieved by implementing the `_fit_resample` abstract method of the parent class `BaseOverSampler`. More specifically, the function `_make_geometric_sample` implements the data generation mechanism of G-SMOTE as shortly described in section 1.3. This function is called in the `_make_geometric_samples` method of the `GeometricSMOTE` class in order to generate the appropriate number of synthetic data for a particular minority class. Finally, the method `_make_geometric_samples` is called in `_fit_resample` method to generate synthetic data for all minority classes. Figure 3 provides a visual representation of the above classes and functions hierarchy.

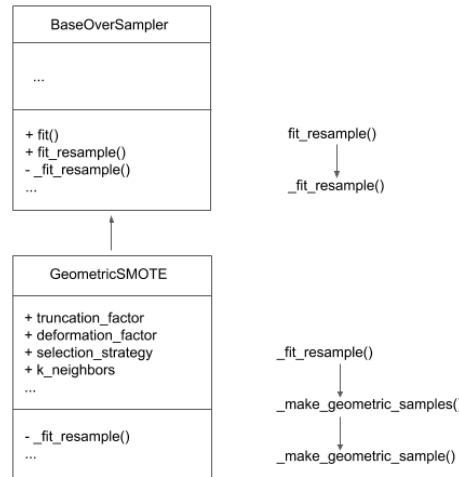


Figure 3: UML class diagrams and callgraphs of main classes and methods.

125 2.2. Software Functionalities

126 As it was mentioned in subsection 2.1, the class `GeometricSMOTE` repre-
127 sents the G-SMOTE oversampler. The initializer of `GeometricSMOTE` includes
128 the following G-SMOTE’s hyperparameters: `truncation_factor`, `deformation_factor`,
129 `selection_strategy` and `k_neighbors` as explained in subsection ???. Once
130 the `GeometricSMOTE` object is initialized with a specific parametrization, it
131 can be used to resample the imbalanced data represented by the input ma-
132 trix `X` and the target labels `y`. Following the Scikit-Learn API, both `X`, `y` are
133 array-like objects of appropriate shape.

134 Resampling is achieved by using the two main methods of `fit` and `fit_resample`
135 of the `GeometricSMOTE` object. More specifically, both of them take as in-
136 put parameters the `X` and `y`. The first method computes various statistics
137 which are used to resample `X` while the second method does the same but
138 additionally returns a resampled version of `X` and `y`.

139 The `geometric-smote` project has been designed to integrate with the
140 Imbalanced-Learn toolbox and Scikit-Learn ecosystem. Therefore the `GeometricSMOTE`
141 object can be used in a machine learning pipeline, through Imbalanced-
142 Learn’s class `Pipeline`, that automatically combines `samplers`, `transformers`
143 and `estimators`. The next section provides examples of the above function-
144 alities.

145 3. Illustrative Examples

146 3.1. Basic example

147 An example of resampling multi-class imbalanced data using the `fit_resample`
148 method is presented in Listing 1. Initially, a 3-class imbalanced dataset is
149 generated. Next, `GeometricSMOTE` object is initialized with default values for
150 the hyperparameters, i.e. `truncation_factor = 1.0`, `deformation_factor =`
151 `0.0`, `selection_strategy = combined`. Finally, the object’s `fit_resample`
152 method is used to resample the data. Printing the class distribution before
153 and after resampling confirms that the resampled data `X_res`, `y_res` are per-
154 fectly balanced. `X_res`, `y_res` can be used as training data for any classifier
155 in the place of `X`, `y`.

Listing 1: Resampling of imbalanced data using the `fit_resample` method.

```
156 # Import classes and functions.  
157 from collections import Counter  
158 from gsmote import GeometricSMOTE  
159 from sklearn.datasets import make_classification  
160  
161 # Generate an imbalanced 3-class dataset.
```

```

162 X, y = make_classification(
163     random_state=23,
164     n_classes=3,
165     n_informative=5,
166     n_samples=500,
167     weights=[0.8, 0.15, 0.05]
168 )
169
170 # Create a GeometricSMOTE object with default hyperparameters.
171 gsmote = GeometricSMOTE(random_state=10)
172
173 # Resample the imbalanced dataset.
174 X_res, y_res = gsmote.fit_resample(X, y)
175
176 # Print number of samples per class for initial and resampled data.
177 init_count = list(Counter(y).values())
178 resampled_count = list(Counter(y_res).values())
179
180 print(f'Initial_class_distribution: {init_count}.')
181 # Initial class distribution: [400, 75, 25].
182
183 print(f'Resampled_class_distribution: {resampled_count}.')
184 # Resampled class distribution: [400, 400, 400].

```

185 3.2. Machine learning pipeline

186 As mentioned before, the **GeometricSMOTE** object can be used as a part
187 of a machine learning pipeline. Listing 2 presents a pipeline composed by a
188 G-SMOTE oversampler, a PCA tranformation and a decision tree classifier.
189 The pipeline is trained on imbalanced binary-class data and evaluated on a
190 hold-out set. The user applies the process in a simple way while the internal
191 details of the calculations are hidden.

Listing 2: Training and evaluation of a machine learning pipeline that contains the **GeometricSMOTE** object.

```

192 # Import classes and functions.
193 from gsmote import GeometricSMOTE
194 from sklearn.datasets import make_classification
195 from sklearn.decomposition import PCA
196 from sklearn.tree import DecisionTreeClassifier
197 from sklearn.model_selection import train_test_split
198 from sklearn.metrics import f1_score

```



```

199 from imblearn.pipeline import make_pipeline
200
201 # Generate an imbalanced binary-class dataset.
202 X, y = make_classification(
203     random_state=23,
204     n_classes=2,
205     n_samples=500,
206     weights=[0.8, 0.2]
207 )
208
209 # Split the data to training and hold-out sets.
210 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=
211
212 # Create the pipeline's objects with default hyperparameters.
213 gsmote = GeometricSMOTE(random_state=11)
214 pca = PCA()
215 clf = DecisionTreeClassifier(random_state=3)
216
217 # Create the pipeline.
218 pip = make_pipeline(gsmote, pca, clf)
219
220 # Fit the pipeline to the training set.
221 pip.fit(X_train, y_train)
222
223 # Evaluate the pipeline on the hold-out set using the F-score.
224 test_score = f1_score(y_test, pip.predict(X_test))
225
226 print(f'F-score on hold-out set: {test_score}.')
227 # F-score on hold-out set: 0.7.

```

228 4. Impact and conclusions

229 Classification of imbalanced datasets is a challenging task for standard
230 machine learning algorithms. G-SMOTE, as an enhancement of the SMOTE
231 data generation mechanism, provides a flexible and effective way for resam-
232 pling the imbalanced data. G-SMOTE's empirical results prove that it out-
233 performs SMOTE and its variants. Machine learning researchers and indus-
234 try practitioners can benefit from using G-SMOTE in their work since the
235 imbalanced learning problem is a common characteristic of many real-world
236 applications.

237 The `geometric-smote` project provides the only Python implementation,
 238 to the best of our knowledge, of the state-of-the-art oversampling algorithm
 239 G-SMOTE. A significant advantage of this implementation is that it is built
 240 on top of the Scikit-Learn’s ecosystem. Therefore, using the G-SMOTE
 241 oversampler in typical machine learning workflows is an effortless task for the
 242 user. Also, the public API of the main class `GeometricSMOTE` is identical to
 243 the one implemented in Imbalanced-Learn for all oversamplers. This means
 244 that users of Imbalanced-Learn and Scikit-Learn, that apply oversampling
 245 on imbalanced data, can integrate the `gsmote` package in their existing work
 246 in a straightforward manner or even replace directly any Imbalanced-Learn’s
 247 oversampler with `GeometricSMOTE`.

248 5. Conflict of Interest

249 We wish to confirm that there are no known conflicts of interest associated
 250 with this publication and there has been no significant financial support for
 251 this work that could have influenced its outcome.

252 References

- 253 [1] N. V. Chawla, A. Lazarevic, L. Hall, K. Boyer, SMOTEBoost: im-
 254 proving prediction of the minority class in boosting, *Principles of*
 255 *Knowledge Discovery in Databases, PKDD-2003* (2003) 107–119doi:
 256 10.1007/b13634.
 257 URL [http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.](http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.80.1499)
 258 [1.80.1499](http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.80.1499)
- 259 [2] G. Haixiang, L. Yijing, J. Shang, G. Mingyun, H. Yuanyue, G. Bing,
 260 Learning from class-imbalanced data: Review of methods and appli-
 261 cations, *Expert Systems with Applications* 73 (2017) 220–239. doi:
 262 10.1016/j.eswa.2016.12.035.
 263 URL <https://doi.org/10.1016/j.eswa.2016.12.035>
- 264 [3] N. V. Chawla, K. W. Bowyer, L. O. Hall, W. P. Kegelmeyer, SMOTE:
 265 Synthetic minority over-sampling technique, *Journal of Artificial Intel-*
 266 *ligence Research* 16 (2002) 321–357. arXiv:1106.1813, doi:10.1613/
 267 jair.953.
- 268 [4] S. Barua, M. M. Islam, X. Yao, K. Murase, MWMOTE - Majority
 269 weighted minority oversampling technique for imbalanced data set learn-
 270 ing, *IEEE Transactions on Knowledge and Data Engineering* 26 (2)
 271 (2014) 405–425. doi:10.1109/TKDE.2012.232.

- [5] A. Fernández, V. López, M. Galar, M. J. del Jesus, F. Herrera, Analysing the classification of imbalanced data-sets with multiple classes: Binarization techniques and ad-hoc approaches, *Knowledge-Based Systems* 42 (2013) 97–110. doi:<http://dx.doi.org/10.1016/j.knosys.2013.01.018>. URL <http://www.sciencedirect.com/science/article/pii/S0950705113000300>
- [6] H. He, E. A. Garcia, Learning from Imbalanced Data, *IEEE Transactions on Knowledge and Data Engineering* 21 (9) (2009) 1263–1284. arXiv:arXiv:1011.1669v3, doi:10.1109/TKDE.2008.239.
- [7] G. Lemaitre, F. Nogueira, C. K. Aridas, Imbalanced-learn: A Python Toolbox to Tackle the Curse of Imbalanced Datasets in Machine Learning, *Journal of Machine Learning Research* 18 (2016) 1–5. arXiv:1609.06570, doi:<http://www.jmlr.org/papers/volume18/16-365/16-365.pdf>. URL <http://arxiv.org/abs/1609.06570>
- [8] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, E. Duchesnay, Scikit-learn: Machine learning in Python, Vol. 12, 2011. arXiv:arXiv:1201.0490v2, doi:10.1007/s13398-014-0173-7.2. URL <http://dl.acm.org/citation.cfm?id=2078195>
- [9] R. C. Prati, G. Batista, M. C. Monard, Learning with class skews and small disjuncts, in: *SBIA*, 2004, pp. 296–306. doi:10.1007/978-3-540-28645-5_30.
- [10] G. Douzas, F. Bacao, Self-organizing map oversampling (somo) for imbalanced data set learning, *Expert Systems with Applications* 82 (2017) 40 – 52. doi:<https://doi.org/10.1016/j.eswa.2017.03.073>. URL <http://www.sciencedirect.com/science/article/pii/S0957417417302324>
- [11] G. Douzas, F. Bacao, F. Last, Improving imbalanced learning through a heuristic oversampling method based on k-means and smote, *Information Sciences* 465 (2018) 1 – 20. doi:<https://doi.org/10.1016/j.ins.2018.06.056>. URL <http://www.sciencedirect.com/science/article/pii/S0020025518304997>